

Stateless Best Effort Multicast

draft-lx-msr6-rgb-segment-03
draft-chen-pim-be-mrh-00

Yisong Liu (China Mobile)
Jingrong Xie, Xuesong Geng (China Telecom)
Mengxiao Chen (New H3C Technologies)

Huaimo Chen, Donald E. Eastlake, Mike McBride (Futurewei)
Yanhe Fan (Casa Systems)
Gyan Mishra (Verizon)
Yisong Liu (China Mobile)
Aijun Wang (China Telecom)
Xufeng Liu (IBM Corporation)
Lei Liu (Fujitsu)

MSR6 BoF IETF 114

Brief Description

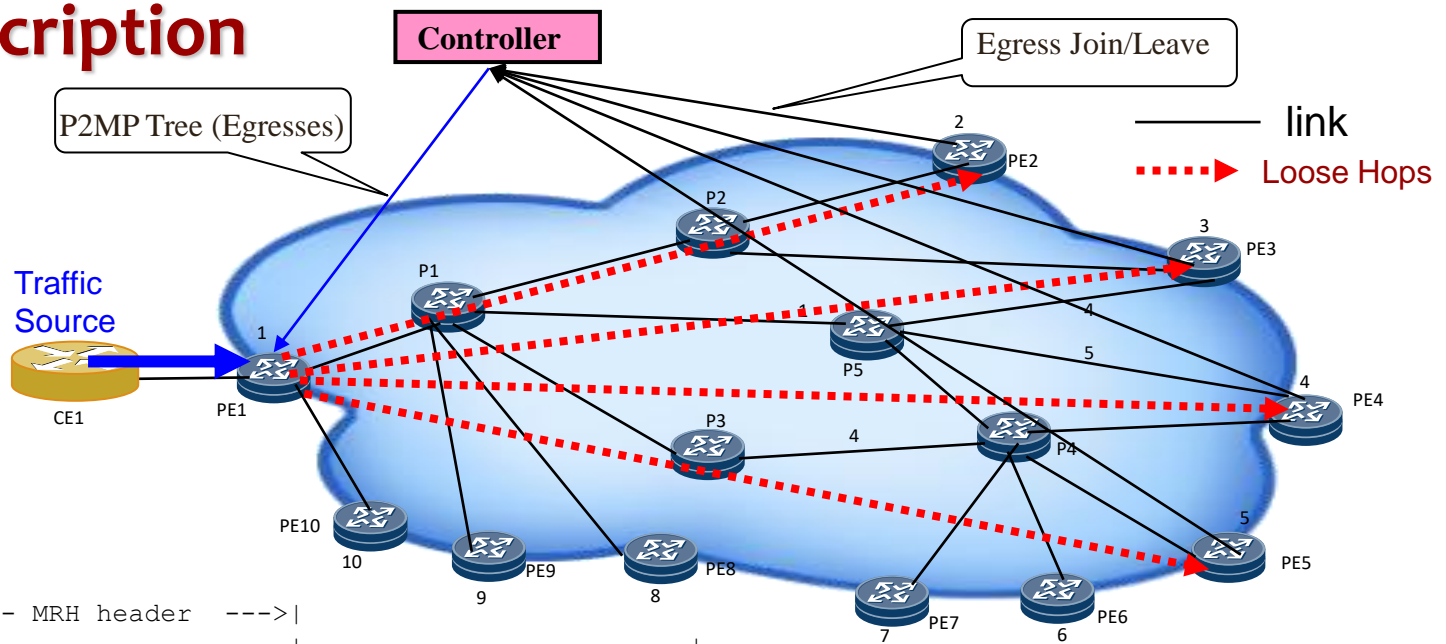


Figure 1. Tree from PE1 towards PE2 to PE5

```
|<---IPv6 header-->|<--- MRH header  --->|
+-----+-----+-----+
| Next Header =    | Next Header          | (an extension header) |
| 60 (DoH) or 43 (RT) |                  | IP multicast datagram |
| SA=IPv6 Address  | Type = TBD (MRH)      |                      |
| DA=IPv6 Address  | Data (Subtree/Egresses) |                      |
+-----+-----+-----+
|<----- MRH ----->|
+-----+-----+-----+
| Next Header      | Hdr Ext Len | Type = TBD | . . . |
+-----+-----+-----+
|                  |                  |                  |                  |
+=====+
|                  | Tree/Sub-tree (i.e., egresses) encoded |
|                  |                  |                  |                  |
+=====+
```

- Ingress (e.g., PE1) encapsulates the packet in a MRH with tree (i.e., egresses of tree)
- The packet is transmitted along the shortest IGP pathes to the egresses.
- Egress (e.g., PE2) decapsulates the packet in a MRH and sends it to next header process

Solution 1 Specials: DoH

|<--IPv6 header-->|<-- DoH header -->|

Next Header =	Next Header	(an extension header)
60 (DoH header)		IP multicast datagram
SA=IPv6 Address	Option Type =	
DA=IPv6 Address	TBD (MRH)	
	Data (Sub-tree)	

|<---- MRH ---->|

Next Header	Hdr Ext Len	Option Type=TBD	Option Length
BIFT-id		Rsv	TTL
Rsv	Ver	BSL	Entropy
OAM	Rsv	DSCP	Rsv
BitString (first 32 bits)			~
~			~
BitString (last 32 bits)			

RGB
Option } IPv6
Data } **MRH** (DoH)
for BE

Solution 1 Specials: Forwarding Procedure

1. IF (There is DoH as an IPv6 Extension header and one of the options type is RGB
2. Lookup BIFT based on the bitstring inside the RGB Option Data
3. Forward the packet via the matched entry in the BIFT
4. ELSE IF NH=ICMPv6 or (NH=RGB Extension Header Type and NH of Extension Header=ICMPv6)
5. Send to CPU.
6. ELSE ;Ref
7. Drop the packet.

Ref: An ICMPv6 packet using End.RGB as destination address.

Solution 2 Specials: Routing Header

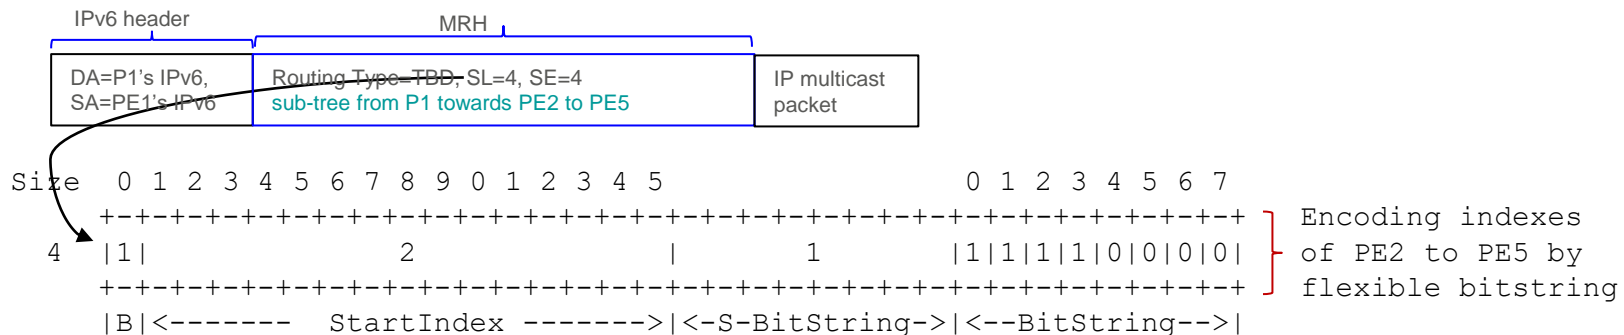
```
|<--IPv6 header-->|<-Routing header->|
```

Next Header =	Next Header	(an extension header)
43(Routing header)		IP multicast datagram
SA=IPv6 Address	Routing Type =	
DA=IPv6 Address	TBD (MRH)	
	SL, SE, Sub-tree	

MRH

[illegible]

- IPv6
 - MRH** (Routing header) for BE



IPv6 packet with one bitstring sent to P1

Solution 2 Specials: Flexible Bitstring and NodeIndex

A flexible bitstring has four fields:

- 1). B flag with value 1, 2). start index (StartIndex), 3). size of bitstring (S-BitString) in bytes and
- 4). bitstring (BitString), where each bit with value 1 indicates a node index equal to StartIndex plus the bit number. Note that the bit number is counted from right to left and from 0.

For example, the indexes of egresses PE2 to PE5 (i.e., PE2, PE3, PE4 and PE5) are encoded by a flexible bitstring (suppose their indexes are 2, 3, 4 and 5 respectively):

B = 1, StartIndex = 2, S-BitString = 1, BitString = 0b11110000 indicating four node indexes 2, 3, 4 and 5.

BitString's first bit (bit 0) with value 1 indicates the first node index 2 equal to 2 + 0; the BitString's second bit (bit 1) with value 1 indicates the second node index 3 equal to 2 + 1, and so on.

```

Size  0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5          0 1 2 3 4 5 6 7
      +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
4      |1|                2                |      1      |1|1|1|1|0|0|0|0|
      +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
      |B|<----- StartIndex ----->|<-S-BitString->|<--BitString-->|
  
```

Encoding indexes of PE2 to PE5 by flexible bitstring

A NodeIndex field with B = 0 represents a node index directly.

For example, the indexes of egresses PE2 to PE5 are represented by NodeIndex.

```

Size  0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5
      +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
8      |0| NodeIndex = 2 (PE2's Index) |
      +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
6      |0| NodeIndex = 3 (PE3's Index) |
      +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
4      |0| NodeIndex = 4 (PE4's Index) |
      +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
2      |0| NodeIndex = 5 (PE5's Index) |
      +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
      |B|<----- NodeIndex ----->|
  
```

Encoding indexes of PE2 to PE5 by NodeIndex

Indexes of PE2 to PE5 by string:
4 bytes, operations on 8 bits

Indexes of PE2 to PE5 by NodeIndex:
8 bytes, operations on 4 Node Indexes

Using bitstring is more efficient
than using NodeIndex

Solution 2 Specials: More Efficient Encoding

For a tree (i.e., egress nodes of tree), optimal or more efficient encoding is selected and used in MRH wrt space and processing time.

For example, for a tree from PE1 to PE2 – PE5,

if PE1 – PE5 have their indexes 2 – 5 respectively, encoding by flexible bitstring is selected;

If PE1 – PE5 have their indexes 102, 503, 1004 and 1005 respectively, encoding by NodeIndex and flexible bitstring is selected.

Indexes of PE2 and PE3 by NodeIndex:
4 bytes, operations on 2 Node Indexes

```
Size  0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5
      +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
      8 |0|NodeIndex = 102 (PE2's Index)|
      +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
      6 |0|NodeIndex = 503 (PE3's Index)|
      +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
      4 |1|                               1004                               |
      +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
      |B|<----- StartIndex ----->|<-S-BitString->|<--BitString-->|
```

Indexes of PE2 and PE3 by NodeIndex

Indexes of PE2 – PE5 by NodeIndex and flexible bitstring

Indexes of PE2 and PE3 by string:
54 bytes, operations on 408 bits

Using NodeIndex and bitstring
is more efficient than using 2
bitstrings

```
Size  0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5
      +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
      58 |1|                               102                               |
      +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
      4 |1|                               1004                               |
      +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
      |B|<----- StartIndex ----->|<-S-BitString->|<--BitString-->|
```

Indexes of PE2 and PE3 by flex bitstring

Indexes of PE4 and PE5 by flexible bitstring

Next

Comments